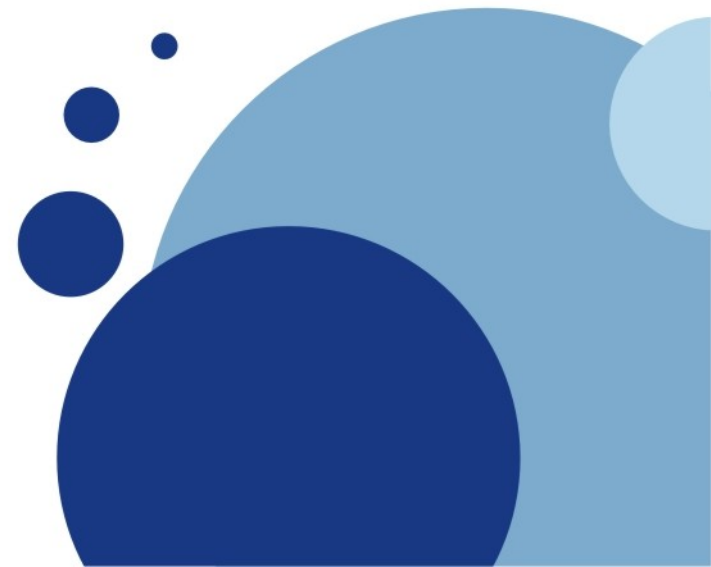


Materialized Views

Dr. Rim Moussa

University of Carthage
`rim.moussa@gmail.com`



Aggregate Tables

- Aggregate tables are also known as :
 - Snapshots, summary table, materialized views (Oracle), materialized query tables (DB2), indexed views (SQL Server)



- Definition


- *A materialized view* is a table that contains the result of a query
- if a view is used frequently enough, it may even be efficient to materialize it
- Aggregates are pre-calculated summaries derived from the most granular fact table
- Goals: Aggregate table are used to:
 - Cache expensive queries in a data warehouse with summarized table. Then use the cache to process queries.
 - Replicate data to non-master sites in a replication environment
- MVs are recommended when queries are known:
 - OLAP queries are usually ad-hoc!

Cost of Aggregate Tables

- Complex calculus
 - optimized through smart data fragmentation and distributed (parallel) calculus
- Storage requirement
 - Volume in bytes of each MV
- Refresh
 - recompute parts of the materialized view each time one of the underlying base tables changes.
 - How frequent are changes?
 - Always compare
 - incremental refresh performances to
 - recomputation for source base tables performances

What to materialize?

- MVs Advisor
 - Exhaustive enumeration
 - Recommendations for a specific workload
 - m number of dimensions, number of nodes in the lattice = 2^m
- Data cube lattice: 4 dimensions: product,store,day,reduction



q_1 = total sales per product

q_2 = total sales per product and store

q_3 = total sales per product and day

q_4 = total sales per product, store and day

Example Q12 of TPC-H benchmark

```
SELECT l_shipmode,  
sum(case when o_orderpriority = '1-URGENT' OR o_orderpriority  
= '2-HIGH' then 1 else 0 end) as high_line_count,  
sum(case when o_orderpriority <> '1-URGENT' AND  
o_orderpriority <> '2-HIGH' then 1 else 0 end) AS  
low_line_count  
  
FROM  orders, lineitem  
  
WHERE  
    o_orderkey = l_orderkey  
    AND l_shipmode in ('MAIL', 'SHIP')  
    AND l_commitdate < l_receiptdate  
    AND l_shipdate < l_commitdate  
    AND l_receiptdate >= date '1994-01-01'  
    AND l_receiptdate < date '1994-01-01' + interval '1' year  
  
GROUP BY  l_shipmode  
ORDER BY  l_shipmode;
```

Parameters:

_list of values for *l-shipmode*

_1st day of a given year

MV12 of TPC-H benchmark (CTAS statement)

```
CREATE TABLE agg_c12 AS
```

```
SELECT year, l_shipmode,  
sum(case when o_orderpriority = '1-URGENT' or o_orderpriority  
='2-HIGH' then 1 else 0 end) as high_line_count,  
sum(case when o_orderpriority <> '1-URGENT' and  
o_orderpriority <> '2-HIGH' then 1 else 0 end) as  
low_line_count  
FROM orders, lineitem, time
```

```
WHERE o_orderkey = l_orderkey  
AND timekey_receiptdate = time.timekey  
AND o_orderkey = l_orderkey  
AND l_commitdate < l_receiptdate  
AND l_shipdate < l_commitdate
```

```
GROUP BY year, l_shipmode  
ORDER BY year, l_shipmode;
```

Calculate agg_c12 metadata:

_volume (bytes)

_cardinality: DW age (number of years) X |l_shipmode|

_time to build (sec)

Q12 re-written

```
SELECT l_shipmode, high_line_count, low_line_count  
  
FROM   agg_c12  
  
WHERE  l_shipmode in ('MAIL', 'SHIP')  AND year = '1994'  
  
ORDER BY  l_shipmode;
```

Gain in performance because we don't

- _ filter *lineitem* (*l_commitdate* < *l_receiptdate* AND *l_shipdate* < *l_commitdate*) : full table scan
- Lineitem* has 6M X SF records
- SF is the TPC-H scale factor
- SF = 1 --> TPC-H warehouse volume is 1GB of raw data
- _ run expensive joins
- orders* ⋈ *lineitem* is not performed
- Lineitem* has 6M X SF records
- Orders* has 1.5M X SF records

And measures are pre-computed

Materialized Views Data for TPC-H benchmark

MV-Qi	Volume (MB)
mv-q1	0.008
mv-q3	52.712
mv-q4	2.241
mv-q5	2.563
mv-q6	0.088
mv-q7	0.067
mv-q8	2.128
mv-q12	0.002
mv-q13	0.003
mv-q14	0.001
mv-q15	0.001
mv-q16	2.861
mv-q17	0.011
mv-q18	0.023
mv-q19	836.278
mv-q22	7.630
	<i>901.817</i>

--MonetDB

```
SELECT table, sum(columnsize)  
FROM storage()  
GROUP BY table;
```

Less than 1GB for MVs whether is the scale factor!

Materialized Views Refresh

RF1: new inserts in ORDERS and LINEITEM


```
-- create a temp lineitem table
CREATE TABLE LINEITEM_TEMP (L_ORDERKEY INTEGER NOT NULL,
L_PARTKEY .. L_PROFIT DECIMAL(15,2));
-- create a temp orders table
CREATE TABLE ORDERS_TEMP (O_ORDERKEY INTEGER NOT NULL, ...
O_SUM_LOST_REVENUE DECIMAL(15,2));
--load data into temp table
--Oracle
sqlldr system/manager@XE control = ins_orders.ctl
sqlldr system/manager@XE control = ins_lines.ctl
.ctl
load data
infile '/home/oracle/TPCCDATA/orders_*.tbl'
into table orders_tmp
fields terminated by "|"
(O_ORDERKEY, ...)
--MonetDB
COPY INTO LINEITEM_TEMP FROM '$HOME\refresh\ins_lines.tbl'
USING DELIMITERS '|', '\n' ;
COPY INTO ORDERS_TEMP FROM '$HOME\refresh\ins_orders.tbl'
USING DELIMITERS '|', '\n' ;
```

Incremental MVs Refresh (PL or Merge Statement)

```
CREATE TABLE agg_c12_ins AS
SELECT time.year as year, l_shipmode,
sum(case when o_orderpriority = '1-URGENT' or o_orderpriority = '2-HIGH' then 1
else 0 end) as high_line_count,
sum(case when o_orderpriority <> '1-URGENT' and o_orderpriority <> '2-HIGH'
then 1 else 0 end) as low_line_count
FROM orders_temp, lineitem_temp, time
WHERE timekey_receiptdate = time.timekey AND o_orderkey = l_orderkey
AND l_commitdate < l_receiptdate AND l_shipdate < l_commitdate
GROUP BY time.year, l_shipmode;
```

```
DECLARE
CURSOR delta IS SELECT * FROM agg_c12_ins;
BEGIN
FOR d IN delta LOOP
    UPDATE agg_c12 SET
        high_line_count = high_line_count + d.high_line_count,
        low_line_count = low_line_count + d.low_line_count,
        WHERE year = d.year and l_shipmode = d.l_shipmode;
END LOOP;
END;
/
drop table agg_c12_ins;
COMMIT;
```

Oracle Materialized Views: scenario



```
CREATE MATERIALIZED VIEW view-name  
BUILD [IMMEDIATE | DEFERRED]  
REFRESH [FAST | COMPLETE | FORCE ]  
ON [COMMIT | DEMAND ]  
[[ENABLE | DISABLE] QUERY REWRITE]  
AS  
SELECT ...;
```

```
CREATE MATERIALIZED VIEW LOG ON scott.emp  
TABLESPACE users  
WITH PRIMARY KEY  
INCLUDING NEW VALUES;
```

Oracle Materialized Views Creation

- BUILD

- IMMEDIATE: the MV is populated immediately.
- DEFERRED: The MV is populated on the first requested refresh.

- REFRESH

- FAST: a fast refresh is attempted
 - Need to create an MV log
 - Incremental refresh only for simple SQL
- COMPLETE: The table segment supporting the MV is truncated and repopulated completely using the associated query
- FORCE: A fast refresh is attempted else a complete refresh is performed.
- COMMIT: The refresh is triggered by a committed data change in one of the dependent tables.
- DEMAND: The refresh is initiated by a manual request or a scheduled task.

```
EXEC DBMS_MVIEW.refresh('EMP_MV');
```

Oracle Materialized Views

--privileges to create an MV

CREATE MATERIALIZED VIEW

CREATE TABLE --an MV is a table

GRANT SELECT --on tables required for building the MV

--enable QUERY REWRITE

GRANT QUERY REWRITE TO user|role;

ALTER SESSION SET

QUERY_REWRITE_ENABLED = TRUE;

--data about refresh jobs

SELECT job, SCHEMA_USER,

TO_CHAR(last_date,'DD/MM/YYYY HH:MI') "last refresh",

TO_CHAR(next_date,'DD/MM/YYYY HH:MI') "next refresh",

interval "when ",

what "what"

FROM dba_jobs

WHERE what LIKE '%refresh%' AND SCHEMA_USER='SCOTT';

Big Picture

