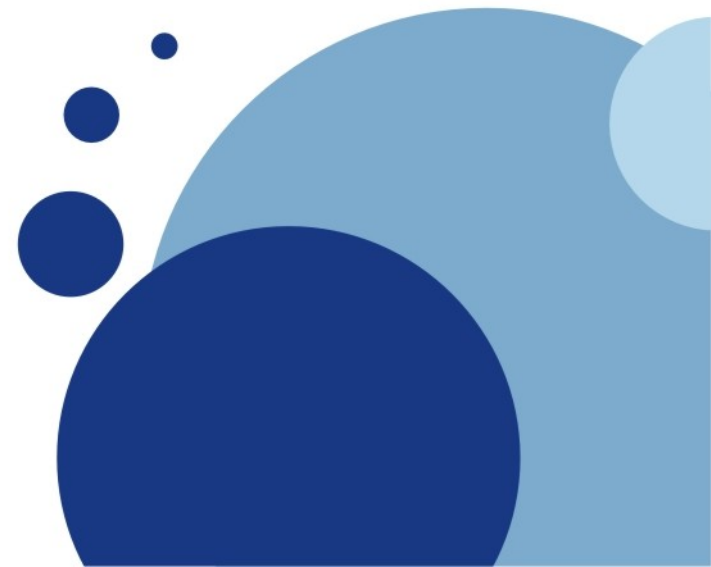


# OLAP Indexes

**Dr. Rim Moussa**

University of Carthage  
`rim.moussa@gmail.com`



# Indexes

- Indexes
  - B-trees, hash tables,
  - Grids, quad-trees, R-trees (spatial data)
- OLAP indexes
  - Bitmaps, inverted lists,
  - Join Indexes, Bitmap join indexes
  - Column store indexes
- ROWID
  - physical address of a row
  - Oracle
    - OOOOOOO.FFF.BBBBBBB.RRR

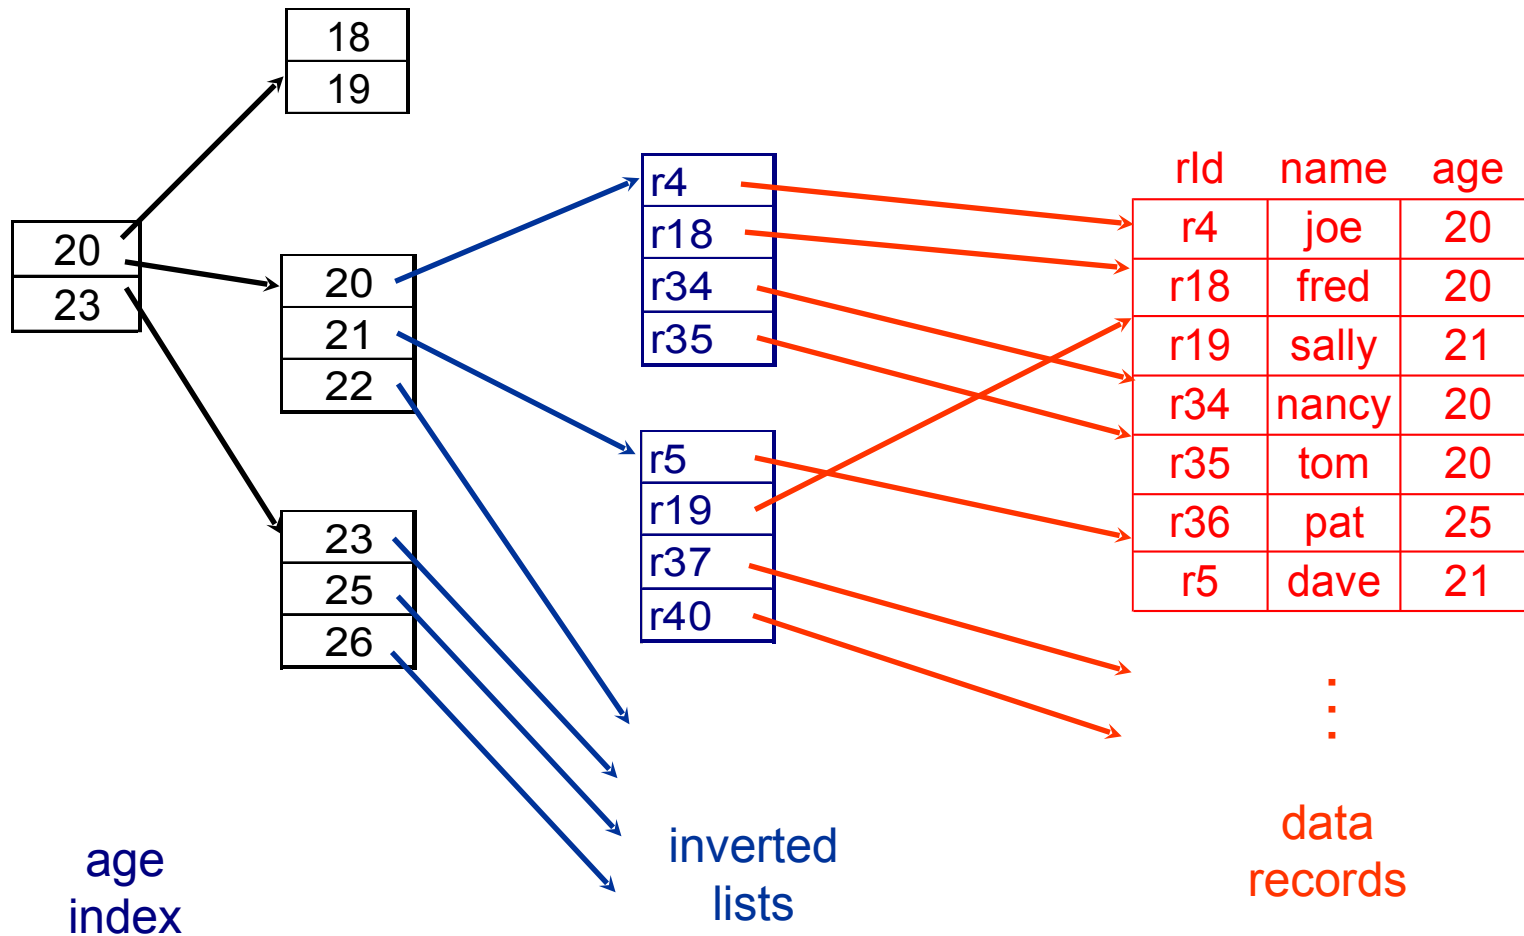
OOOOOO is the object ID

FFF is the file number

BBBBBB is the block number

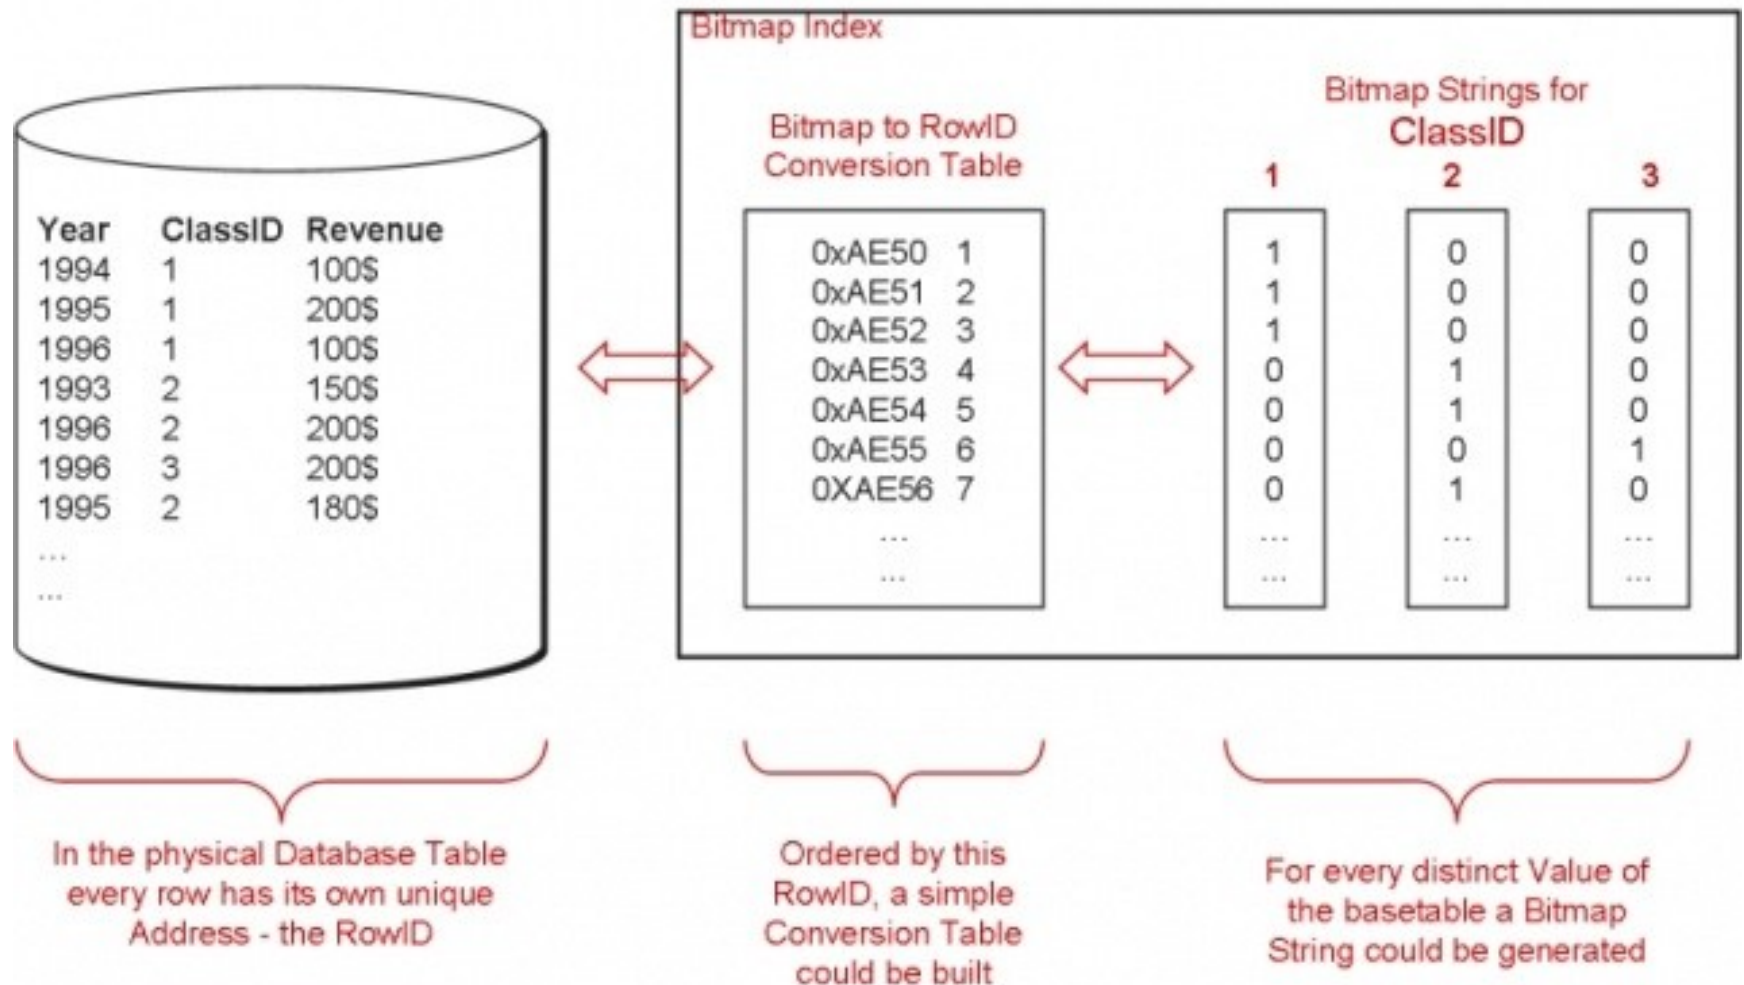
RRR is the row number

# Inverted List Index



By Hector Garcia Molina: Data Warehousing and OLAP

# Bitmap Index



# Operations on Bit Vectors

Gender	MaritalStatus	ChildrenYN	Income	HomeOwner
M	Married	N	120000	N
M	Single	N	80000	Y
F	Divorced	Y	75000	N
F	Married	Y	70000	Y
F	Married	Y	130000	Y
M	Single	Y	45000	N
F	Married	N	42000	N

Male	Female
1	0
1	0
0	1
0	1
0	1
1	0
0	1

Married	Single	Divorced
1	0	0
0	1	0
0	0	1
1	0	0
1	0	0
0	1	0
1	0	0

Y	N
0	1
1	0
0	1
1	0
1	0
0	1
0	1

SELECT \* FROM survey WHERE Gender='Male' AND MaritalStatus='Single' AND HomeOwner='Y'

Male		Single		Y			
1		0		0		0	
1		1		1		1	
0		0		0		0	
0	AND	0		1		0	
0		0	AND	1		0	
1		1		0		0	
0		0		0		0	

		EQUALS	
0			0
1			1
0			0
0			0
0			0
0			0


This row satisfies the query

# Join Index

join index

product	id	name	price	jIndex
	p1	bolt	10	r1,r3,r5,r6
	p2	nut	5	r2,r4

sale	rld	prodld	storeld	date	amt
	r1	p1	c1	1	12
	r2	p2	c1	1	11
	r3	p1	c3	1	50
	r4	p2	c2	1	8
	r5	p1	c1	2	44
	r6	p1	c2	2	4



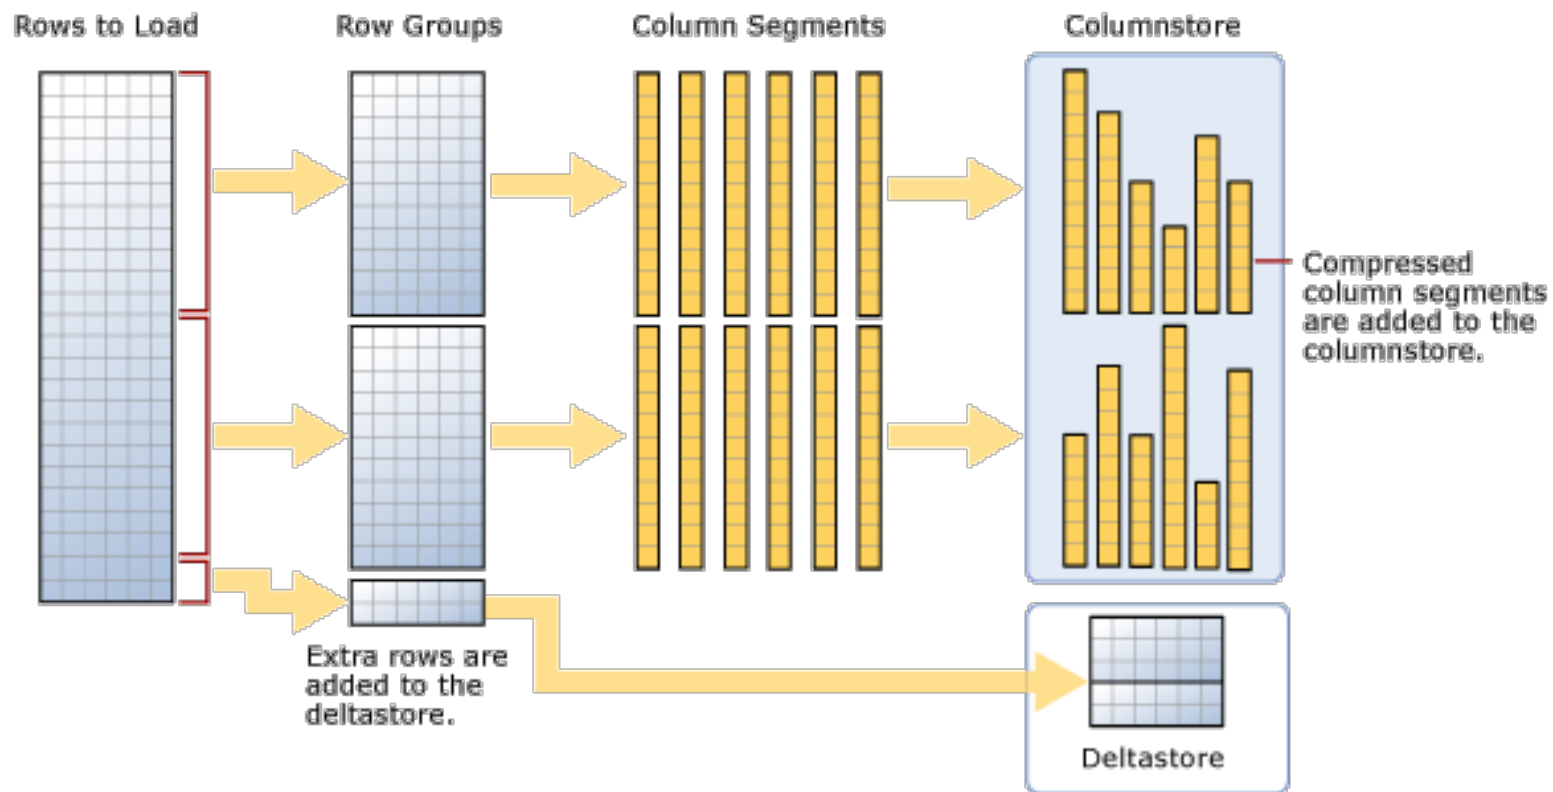
# Bitmap Join Index

Client Table			
RID <sup>1</sup>	CID	Name	Country
6	616	Gilles	France
5	515	Ives	USA
4	414	Patrick	Tunisia
3	313	Odiel	Tunisia
2	212	Eric	France
1	111	Pascal	France

Sales Table				
RID <sup>2</sup>	CID	PID	TID	Price
1	616	106	11	85
2	616	106	66	26
3	616	104	83	50
4	515	104	11	10
5	414	109	66	14
6	212	106	55	14
7	111	103	44	20
8	111	103	33	27
9	212	103	11	100
10	313	103	11	200
11	414	103	11	103
12	414	103	55	108

Bitmap Join Index			
RID	France	USA	Tunisia
1	1	0	0
2	1	0	0
3	1	0	0
4	0	1	0
5	0	0	1
6	1	0	0
7	1	0	0
8	1	0	0
9	1	0	0
10	0	0	1
11	0	0	1
12	0	0	1

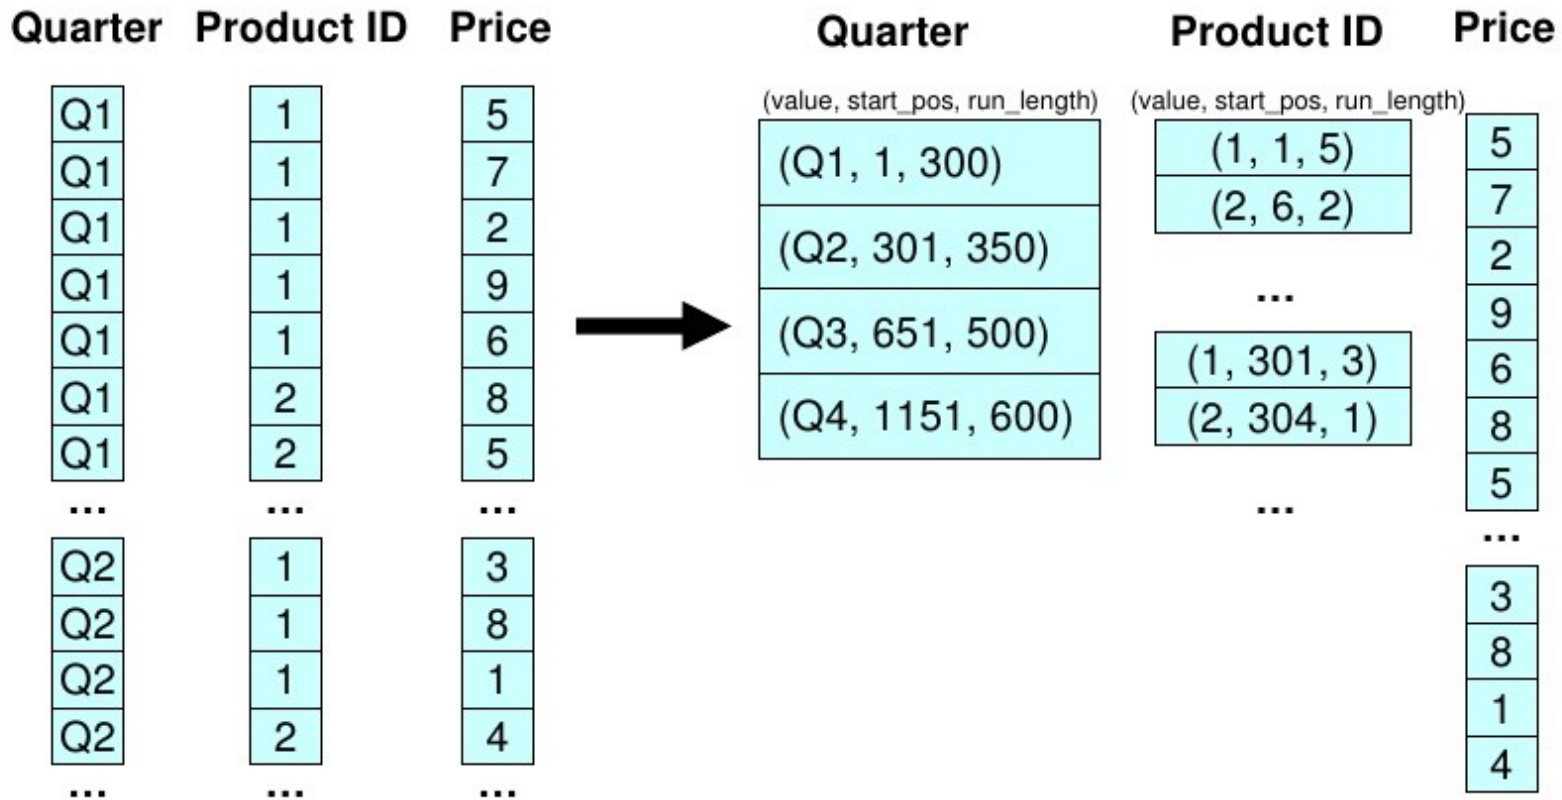
# ColumnStore Index (Microsoft)



- A rowgroup is a group of rows that are compressed into columnstore format at the same time.
- The maximum number of rows per rowgroup is 1,048,576 rows.



# RLE: Run Length Encoding



<https://www.slideshare.net/abadid/vldb-2009-tutorial-on-columnstores>

# Bit-vector Encoding

Re-use permitted when acknowledging the

“Integrating Compression and Execution in Column-Oriented Database Systems” Abadi et. al, SIGMOD '06



## Bit-vector Encoding

- 1 For each unique value,  $v$ , in column  $c$ , create bit-vector  $b$

- 1  $b[i] = 1$  if  $c[i] = v$

- 1 Good for columns with few unique values

- 1 Each bit-vector can be further compressed if sparse

Product ID

1
1
1
1
1
2
2

...

1
1
2
3

...



ID: 1

1
1
1
1
1
0
0

...

1
1
0
0

...

ID: 2

0
0
0
0
0
1
1

...

0
0
1
0

...

ID: 3

0
0
0
0
0
0
0

...

0
0
0
1

...

...

0
0
0
0
0
0
0

...

0
0
0
0

...

# Dictionary Compression


Re-use permitted when acknowledging the

“Integrating Compression and Execution in Column-Oriented Database Systems” Abadi et. al, SIGMOD '06



## Dictionary Encoding

- 1 For each unique value create dictionary entry
- 1 Dictionary can be per-block or per-column
- 1 Column-stores have the advantage that dictionary entries may encode multiple values at once



# Dictionary Compression vs RLE

## Classical DB

Company [CHAR50]	Region [CHAR30]	Group [CHAR5]
INTEL	USA	A
Siemens	Europe	B
Siemens	Europe	C
SAP	Europe	A
SAP	Europe	A
IBM	USA	A

## NewDB Column Store: Dictionary compressed

0 INTEL 1 Siemens 2 SAP 3 IBM	0 Europe 1 USA	0 A 1 B 2 C
0	1	0
1	0	1
1	0	2
2	0	0
2	0	0
3	1	0

## NewDB Column Store: Run length compressed\*


0 INTEL 1 Siemens 2 SAP 3 IBM	0 Germany 1 USA	0 A 1 B 2 C
1 x „0“	1 x „1“	1 x „0“
2 x „1“	4 x „0“	1 x „1“
2 x „2“	1 x „1“	1 x „2“
1 x „3“		3 x „0“



# Frame of Reference Encoding

- 1 Encodes values as  $b$  bit offset from chosen frame of reference
- 1 Special escape code (e.g. all bits set to 1) indicates a difference larger than can be stored in  $b$  bits
  - 1 After escape code, original (uncompressed) value is written

“Compressing Relations and Indexes” Goldstein, Ramakrishnan, Shaft, ICDE'98



# Differential Encoding


Re-use permitted when acknowledging the original © Stavros Harizopoulos, Daniel Abadi, Peter Boncz (2009)



## Differential Encoding

- 1 Encodes values as  $b$  bit offset from previous value
- 1 Special escape code (just like frame of reference encoding) indicates a difference larger than can be stored in  $b$  bits
  - 1 After escape code, original (uncompressed) value is written
- 1 Performs well on columns containing increasing/decreasing sequences
  - 1 inverted lists
  - 1 timestamps
  - 1 object IDs
  - 1 sorted / clustered columns

“Improved Word-Aligned Binary Compression for Text Indexing”  
Ahn, Moffat, TKDE’06



# What compression scheme to use?

